

Your Scanner Says “Critical.” FedRAMP Asks a Harder Question.

Matthew Venne
Chief Technology Officer, stackArmor

July 2026

*A plain-language companion to “A Deterministic, CVSS–Environmental Method for FedRAMP
VDR/VER Vulnerability Prioritization.”*

The problem in one sentence

A vulnerability scanner tells you how bad a flaw is *in the abstract*; FedRAMP’s new rules ask how bad it is *for the agency whose data sits on that specific box* — and those are not the same question.

A “Critical” CVE on a throwaway sandbox can matter less than a “Low” CVE on the database holding six agencies’ records. Everyone in security knows this intuitively. The hard part is turning that intuition into a number you can defend in an audit, reproduce on every run, and explain to a 3PAO without hand-waving.

That’s the gap this work fills.

What FedRAMP gave us, and what it left out

FedRAMP’s Vulnerability Detection and Response and Vulnerability Evaluation Requirements (VDR/VER) introduce a clean prioritization vocabulary:

- **PAIN** — *Potential Agency Impact*, rated N1 through N5. This is the “how bad for the agency” score.
- A **remediation matrix** — how many days (or hours) you have to fix something, based on its PAIN level, how exploitable it is, and whether it’s reachable from the internet.

What FedRAMP deliberately *didn’t* give us is the missing middle: the function that takes a specific CVE on a specific asset and spits out an N-level. They specified the output buckets and the deadlines, but not the classifier that fills the buckets. Every provider is left to invent that themselves.

An invented, analyst-by-analyst classifier is the worst of both worlds: it’s not reproducible (two analysts disagree), and it’s not defensible (you can’t show your work). So the question becomes: **can we build that classifier in a principled, standardized way instead of making one up?**

The key insight: the classifier already exists

Here’s the move at the heart of the proposal. **The thing FedRAMP asks for** — “re-weight a vulnerability’s impact by what’s at stake on this asset” — is exactly what the CVSS Environmental metrics were built to do.

CVSS (the scoring system behind those “Critical/High/Medium” labels) has a lesser-known *Environmental* group designed for precisely this: it lets the organization consuming a vulnerability re-score it based on how much *their* environment cares about confidentiality, integrity, and availability on the affected system. Those knobs are called CR, IR, and AR — the Confidentiality, Integrity, and Availability *Requirements*.

So the mapping is almost one-to-one:

What FedRAMP asks	What already answers it
How sensitive is this asset’s data?	CVSS Security Requirements (CR / IR / AR)
How big is the potential agency impact?	CVSS Modified Impact Sub-Score
Is it likely to be exploited?	EPSS probability + KEV (known exploited)
Is it reachable from the internet?	Deployment/network context
One agency or many?	A scope multiplier on the result

The payoff: **you don’t need to invent a risk algebra**. You derive what federal consequences are actually in scope, assign each asset its architectural role, evaluate two facts about reachability and exploitation, and the published CVSS math does the rest. The output is deterministic — same inputs, same answer, every analyst, every run.

How it actually works (the no-math version)

Every vulnerability can hurt three things: **secrecy** (confidentiality), **correctness** (integrity), and **uptime** (availability). For each one, you multiply two numbers:

1. *How badly does this flaw break it?* (from the CVE’s own CVSS vector)
2. *How much does this particular system depend on it?* (from the asset’s security-impact profile)

Combine those three into a single damage score from 0 to 1, then translate that score into a plain word: **Minimal, Narrow, Disruptive, or Debilitating**. Add one more fact — does this asset serve one agency or many? — and you land on an N-level from N1 to N5.

That’s the whole engine. The *only* place human judgment enters the core scoring is three cut-points that decide where one word turns into the next — and those are meant to be documented, governed, and back-tested against how your senior analysts actually triage.

Before the asset score: what federal impact is actually possible?

There is one important step before the arithmetic. A cloud service may be capable of serving many agencies and many information types, but a particular deployment has a narrower intended use. NIST SP 800-60 gives provisional confidentiality, integrity, and availability profiles for federal information types. The agency confirms which types it will actually use in the offering and whether any special factors apply.

That produces a three-dimensional **Security Requirements ceiling**. Keep the dimensions. If the system category is High because confidentiality is High while integrity is Moderate and availability is Low, the ceiling is H/M/L — not H/H/H. A generic text box or file upload does not authorize every hypothetical kind of federal data.

This is also where the point of view matters. PAIN asks about impact to the agency, government operations, federal assets, and affected individuals. The CSP may care more about the same event for brand or corporate-reputation reasons, but that separate business risk does not silently raise the agency’s PAIN inputs.

But where does the asset profile come from?

Everything above hinges on that second set of numbers — *how much does this system depend on secrecy, correctness, and uptime?* That’s the uncapped CR/IR/AR **asset security-impact profile**. The profile is the required input; the method does not require one taxonomy for producing it.

A provider can assign the three objectives directly from an evidenced component assessment, derive them through a role-based archetype catalog, or use another governed mapping. The important rules are that CR, IR, and AR can differ, the same evidence produces the same vector, and the record explains why each objective was chosen.

Asset archetypes remain a strong example. Instead of hand-rating every server in isolation, a provider can sort assets into a short catalog of roles based on **what they do inside the service** — a database, a login service, a cache, a message bus, a CI/CD runner, a public load balancer, and so on. Each role maps to a pre-filled and justified profile. Tag an asset with its archetype, and the profile is inherited automatically. The original PAIN paper uses this approach in its worked architectures because it makes the examples concrete; it is not a required CSP taxonomy.

The benefit is the **auditable derivation trail**. A bare “CR: Low, IR: High, AR: Low” assignment can be hard to reconstruct six months later. A *data-backbone message broker* role makes one possible rationale visible and consistent with comparable brokers. A direct assignment can be just as defensible when it retains the architecture, trust, privilege, dependency, data-handling, and consequence evidence behind each objective.

For providers using archetypes, classify by role, not software kind. The same in-memory store can be a throwaway cache, a session-token store, or a job broker; *how it is used* drives the profile, not *what product it is*. The example catalog recommends a “control-plane lens first” rule: if an asset can deploy, orchestrate, hold cross-estate credentials, or actuate configuration, classify it by that powerful control function regardless of the data it happens to store.

A scalar asset-value label is a permitted shortcut: Low becomes L/L/L, Moderate becomes M/M/M, and High becomes H/H/H. It is also deliberately cautioned against as the only interface. It forces all three objectives to move together and records the conclusion with less evidence, so an implementation should always permit an independent CR/IR/AR profile.

So the full pipeline reads top to bottom like this:

THE WHOLE PIPELINE

Derive intended-use ceiling — which federal information types and consequences are in scope → **resolve the asset security-impact profile** — by direct assessment, an archetype, or another governed method → **take the lower value on each objective** — effective CR/IR/AR for confidentiality, integrity, and availability → **feed the CVSS Environmental impact math** — re-weighting the CVE’s raw impact by what’s at stake *here* → **produces the PAIN level (N1–N5)** — plus, with the reachability and exploitation facts, a concrete deadline.

Two important guardrails on any profile derivation:

- **The archetype list is an example, not a standard.** The proposal publishes a sample catalog as an *existence proof*. Every provider may use its own taxonomy, direct objective assessment, or another governed mapping. Agency FIPS-199 and information-type context belongs in the separate ceiling.
- **Mislabeling is a downgrade attack.** Because the whole score rides on the profile, assigning an asset *down* makes its flaws look smaller than they are. Profile assignments, mappings, and overrides should be controlled and reviewed like any other security setting — and unresolved

assets fail loud at H/H/H (see below), never quiet.

The same flaw, two very different verdicts

This is the part worth internalizing, because it's the entire point of the method. Take one identical remote-code-execution CVE:

- **On a crown-jewel datastore whose security-impact profile and intended-use ceiling retain the necessary High requirements, holding multiple agencies' data:** scores Debilitating, multi-agency → **N5**, remediate in hours.
- **On a single-agency throwaway sandbox:** same CVE, scores Disruptive → **N3**, a much longer clock.

The vulnerability didn't change. The *asset* did all the work. And a “Low”-severity availability flaw on a high-uptime message backbone can legitimately land at **N4 or N5** — decoupled entirely from the vendor's qualitative label.

The PAIN level says how *bad*. Two more facts say how *fast*.

The N-level tells you the potential impact, but not your deadline. Two real-world facts tighten the remediation clock:

- **Is it actually being exploited — or likely to be?** This is the *exploitability* axis. A flaw is treated as likely-exploitable if it clears an EPSS probability threshold (EPSS being the statistical “how likely is this to be exploited in the wild” score) **or** if it's known to be actively exploited — for example, listed in CISA's KEV catalog. Either signal is enough. The method also applies a direct-exposure floor: a directly internet-reachable finding with **AV:N/AC:L/PR:N/UI:N** is treated as likely exploitable. Whichever predicate fires moves the finding to a faster remediation column.
- **Is it reachable from the internet?** A vulnerable surface an unauthenticated outside attacker can actually reach — through a public IP, a public load balancer, an ingress, or any edge path that doesn't authenticate first — is far more urgent than the same flaw buried behind authentication on an internal-only service. (A route gated by an identity-aware proxy in a *remote-access* role — your own people, with the application's own login still behind it — counts as *not* internet-reachable: the attacker can't get to the vulnerable code.)

Stack those up and the logic is intuitive: **more severe + more exploitable + more exposed = less time to fix**. A lookup table turns the combination into a concrete deadline, with tighter clocks for higher-assurance providers. So an internet-facing, actively-exploited flaw on a shared multi-agency edge might be a 12-hour fix, while the same-tier flaw that's sub-threshold on EPSS and tucked behind authentication gets a much longer runway.

There's also an optional refinement that can raise the *level* rather than the clock: where an authoritative source reports that exploiting a flaw grants *total* technical impact, the affected dimensions can be floored to High before scoring — pushing a weak-looking CVE up a tier. It never invents impact on a dimension the CVE doesn't touch.

Fail loud, not quiet

A critical design choice: **missing information never makes a problem look smaller**. An asset whose security-impact profile is unresolved defaults to H/H/H — so it scores loudly and surfaces itself for assessment, rather than hiding in the noise. “Unknown” is treated as “serious until proven otherwise.”

Scoring isn't the end — disposition is

A scanner hit isn't a confirmed problem. The flagged code might not be present, might never run, might need a configuration nobody set, or might already be neutralized by a control. Sorting that out is the **disposition** step, and it can only happen *after* a finding is detected.

The method records dispositions using **VEX** (Vulnerability Exploitability eXchange) — a standardized, signable, machine-readable artifact. Crucially, disposition is an *overlay*, *not a re-score*: a finding's computed PAIN is always retained on record even when it's suppressed. That makes every “this doesn't apply to us” decision **auditable, attributable, and reversible** — with stronger evidence required to wave off a high-stakes finding than a trivial internal one.

Why this matters for the bigger picture

FedRAMP's VDR/VER rule is its implementation of CISA's Binding Operational Directive 26-04 — the broader push to move remediation prioritization *off* raw CVSS base scores and *onto* exploitability and mission impact. FedRAMP kept the directive's philosophy but swapped in its own instruments (PAIN in place of SSVC, among others). Because it chose PAIN, an off-the-shelf SSVC tool doesn't satisfy the rule by itself — and the classifier that produces an N-level was exactly the missing piece. This method supplies it in standards-grounded terms.

The bottom line

You can summarize the entire proposal in four claims:

1. **FedRAMP specified the vocabulary and the deadlines but not the classifier.** Everyone has to build one.
2. **You don't have to invent it.** CVSS Environmental metrics already encode “re-weight impact by what's at stake” — which *is* the agency-impact question.
3. **With governed inputs** — an intended-use ceiling, the asset's security-impact profile, and a one-agency-or-many affected scope — every finding gets a deterministic, auditable N-level and a concrete deadline. The inputs change when use, architecture, or tenancy changes; the finding math does not.
4. **It's defensible.** Same inputs, same output, every time — with the full derivation attached to each finding, and a fail-safe that errs toward over-classifying rather than under.

The deeper paper has the worked examples, the exact formulas, the reference architectures, and a sample archetype-to-profile catalog. The catalog is one illustrative derivation method, *not* a required taxonomy or standard. But the idea you came for is simple: **stop asking “how bad is this CVE?” and start asking “how bad is this CVE *here*?”** — and let a formula everyone can check answer it the same way every time.

Read the full technical memo: A Deterministic, CVSS-Environmental Method for FedRAMP VDR/VER Vulnerability Prioritization. Read the upstream derivation: Before the PAIN Equation: Deriving Security-Requirements Ceilings from Intended Federal Information Types.