

PAIN Relief: A Verified-Control Method for Reducing FedRAMP Rev5 VDR/VER PAIN and Exploitability

Matthew Venne
Chief Technology Officer, stackArmor

July 2026

Abstract

The companion memo *A Deterministic, CVSS–Environmental Method for FedRAMP Rev5 VDR/VER Vulnerability Prioritization* derives a finding’s Potential Agency Impact (PAIN) and its remediation deadline. That memo also establishes a mitigation *ladder*: PAIN comes down when cited, signed evidence lowers a Modified impact metric. This paper makes the ladder systematic. It defines **PAIN Relief**: a governed, versioned catalog that pairs a *machine-verified* security control with the specific weakness class (by CWE) it provably counters and the deterministic scoring reduction that pairing earns. Two levers move: verified controls that bound what a successful exploit can do lower a Modified impact metric and step the PAIN level down; verified controls that reduce the probability of successful exploitation lower a *local* estimate of that probability and can move a finding from likely-exploitable to not. Neither lever mutates a published input — not the CVSS base vector, not the EPSS score, not the CISA KEV clock. The method’s guiding line is *enforced-and-machine-verifiable*, not *claimed*: a control earns relief only when enforcement infrastructure applies it and a collector can read it. The catalog is offered as an example; each provider governs its own.

Contents

1	Status and Relationship to the Companion Papers	2
2	The Problem: Good Hygiene Must Pay	2
2.1	The axis that matters: enforced, not claimed	2
3	Two Levers	2
3.1	Impact: lowering a Modified metric	2
3.2	Exploitability: lowering a local estimate	3
4	The Catalog	4
5	Governance	4
6	What PAIN Relief Never Does	5
7	Worked Examples	5
8	Conclusion	5

1 Status and Relationship to the Companion Papers

This document is informational. It does not define a FedRAMP requirement; it specifies a method for lowering the PAIN and exploitability of a finding within the existing VDR/VER rules, building directly on the PAIN memo’s mitigation ladder. It consumes the internet-reachability determination of *An Exclusion-Gated Method for Determining Internet Reachability* and the disposition vocabulary of *A CycloneDX VEX Profile for FedRAMP Rev5 VDR/VER*. The key words MUST, SHOULD, and MAY follow RFC 2119.

2 The Problem: Good Hygiene Must Pay

A provider that runs its workloads with least privilege, deny-by-default egress, read-only filesystems, and enforced runtime protection is measurably harder to harm than one that does not. Yet a raw scan scores an identical CVE identically on both. If secure configuration changes no number a provider must report, it is all cost and no credit, and the rational provider treats hardening as a burden. The PAIN method already refuses that outcome for one input — asset criticality, through the CVSS Environmental requirements — by re-weighting a vulnerability’s impact by what is at stake on the specific asset. PAIN Relief extends the same principle from *what the asset is* to *what protects it*.

IN PLAIN TERMS

A scanner scores a flaw as if every system running it were equally exposed. Two providers run the same flaw; one has locked the system down and one has not. PAIN Relief is how the locked-down provider’s work shows up in the score — lower impact, or a longer clock — so that doing security well is worth something on paper, not just in principle.

2.1 The axis that matters: enforced, not claimed

The tempting mistake is to reward controls a provider *asserts*. Input sanitization, a WAF signature set, “we validate everything” — these are claims about behavior, unverifiable at assertion time and silently perishable on the next refactor. PAIN Relief admits a control only when it is **machine-verifiable** — read from enforced configuration or telemetry a collector can inspect (a NetworkPolicy, a securityContext, an RDS flag) — or carried by a signed, dated, reusable attestation for structure a platform cannot expose. A control whose effectiveness is a claim stays in the VEX disposition lane of the companion profile, decided per finding by a person. The line is not *a priori* versus *a posteriori*; it is *enforced-and-verifiable* versus *asserted*.

3 Two Levers

PAIN Relief moves exactly two things, and keeps them strictly apart. A control belongs to one lever, chosen by its causal mechanism: a control that bounds *what a successful exploit yields* moves impact; a control that reduces *whether the exploit succeeds* moves exploitability. Because the assignment follows the mechanism, a control never scores on both axes, and nothing double-counts.

3.1 Impact: lowering a Modified metric

The PAIN memo scores a finding through the CVSS Modified Impact Sub-Score, built from the per-dimension impacts *C*, *I*, *A* and the asset’s requirements. A verified control that provably counters

a weakness class lowers the affected Modified impact metric from High to Low before that score is computed:

$$\text{High} \longrightarrow \text{Low}, \quad \text{never None.} \quad (1)$$

The floor is deliberate. A claim of *zero* impact is an applicability claim, and it travels as a VEX disposition (`code_not_present`, `code_not_reachable`), not as relief. The reduced *C*, *I*, or *A* flows through the memo’s existing arithmetic unchanged; the PAIN level and deadline follow. There is one rung, and it does not stack: several controls that each argue a dimension down still lower it by a single step, with every contributing control named in the finding’s evidence.

3.2 Exploitability: lowering a local estimate

FedRAMP’s exploitability axis is binary: a finding is likely-exploitable (LEV) or not (NLEV), and the two select different remediation columns. The memo computes

$$\text{LEV} = \text{KEV} \vee (\text{EPSS} \geq \tau) \vee \text{floor}, \quad \tau = 0.70,$$

where the floor is the FRD-LEV rule that an automated unauthenticated internet exploit is likely-exploitable. PAIN Relief never edits the published EPSS score — a global prediction is not a provider’s to rewrite — just as the impact lever never edits the CVSS base vector. Instead a verified control lowers a *local* estimate of exploitation probability, and the binary threshold does the rest:

$$\text{adjustedEPSS} = \max\left(\text{EPSS} \times \prod_{r \in R} \rho_r, \text{EPSS} \times \phi\right), \quad (2)$$

where R is the set of applicable verified likelihood controls, each carrying a *residual factor* $\rho_r \in (0, 1)$ — the share of exploitation risk it leaves behind — and ϕ is a governed floor on total reduction. LEV is then recomputed with `adjustedEPSS` in place of EPSS.

DEFINITION

ρ_r (*residual factor*): a control that halves exploitation likelihood carries $\rho = 0.5$; one that trims it modestly carries $\rho = 0.85$. Lower is stronger. ϕ (*stacking floor*): the most that stacking may reduce the estimate, e.g. $\phi = 0.5$ caps total reduction at one half. Both are governed, back-tested constants, documented and challengeable like the PAIN cut points.

Three properties of Equation (2) matter. First, **it stacks multiplicatively**: independent controls compound, so layered runtime defense is rewarded — this is the one place defense-in-depth earns more than any single control. Second, **the threshold self-calibrates**: a modest control drops a marginal finding ($0.72 \times 0.85 = 0.61$, now NLEV) but not a near-certain one ($0.99 \times 0.85 = 0.84$, still LEV), which is exactly right — one control should not defuse a vulnerability being exploited everywhere. Third, and absolutely, **KEV is frozen**: a Known Exploited Vulnerability stays LEV regardless of ρ , and its BOD 26-04 clock is untouched. The FRD-LEV floor is a reachability premise, not a probability, so it is defeated only in the binary — by a qualifying edge-authentication boundary — never by a residual factor.

IN PLAIN TERMS

EPSS is a worldwide guess at how likely a flaw is to be attacked. We never change that number. We keep a second, local number that starts at EPSS and drops when we can verify a control that makes the attack less likely to work here. If the local number falls below the line, the finding earns the slower clock. Stacked controls push it lower together — but a flaw that is already being exploited in the wild (KEV) never moves, and the change is always shown beside the original so nothing is hidden.

4 The Catalog

A PAIN Relief entry is a row pairing a control, the CWE class it counters, and the move it earns. Rows are keyed to CWE because the finding carries a CWE and the control’s causal story is a statement about a weakness class — “deny-by-default egress breaks the second-stage fetch of injection and staged-execution flaws” is a claim about **CWE-917**, **CWE-94**, and their kin, not about a single CVE. Where a control bounds what *any* code execution yields regardless of how it began, the row keys to an outcome class (all execution-capable weaknesses, including memory corruption) rather than an enumerated list. A finding with no CWE, or an unrecognized one, earns nothing: relief fails open toward the higher score.

Verified control	Counters (CWE)	Lever	Move
Egress default-deny	917, 94, 98, 829	impact	MC/MI High→Low
Read-only root filesystem	22, 434, 73	impact	MI High→Low
No shell in image	78, 77	impact	MC/MI High→Low
Least-privilege DB grant	89	impact	MC/MI High→Low
Verified HA + rate limit	crash-class DoS	impact	MA High→Low
Blocking-mode EDR/RASP	(any)	likelihood	$\rho = 0.85$
Edge auth (remote-access)	(any)	likelihood	defeats FRD-LEV floor

Table 1: Illustrative PAIN Relief rows. The maintained catalog is proprietary and provider-governed; this sample is an existence proof, not a standard — exactly the posture of the archetype catalog in the companion memo. Every applied credit ships its row identifier, rationale, and the catalog version in the finding’s evidence.

5 Governance

Adding a row grants a discount across the estate, so a row is a security setting and is reviewed like one. Each row **MUST** satisfy:

- **A causal story, not an adjective.** “Egress-deny breaks the staged fetch” qualifies; “EDR makes exploitation harder” does not. The row names the mechanism by which the control interrupts the weakness class.
- **Machine-verifiable or attested, never claimed.** The control is proven from enforced configuration, telemetry, or a signed reusable artifact. This is the clause that keeps input sanitization and WAF signature coverage *out* of the catalog and in the VEX lane: both are bets that a pattern matches every variant an adversary can invent, and Log4Shell’s same-week obfuscations are the standing proof they are not verifiable facts.
- **CWE-keyed and fail-open.** A row applies only to the classes it lists; missing or generic CWEs earn nothing.
- **One rung; never None.** Impact moves High to Low only; zero-impact is a VEX applicability claim. Exploitability moves the binary via Equation (2); nothing here touches reachability or KEV dates.
- **Versioned, pinned, and evidenced.** A score is reproducible only against a named catalog release; the row id and version travel in every evidence line. Refusals are recorded, so the catalog is curated, not accumulated.
- **Back-tested before adoption.** Residual factors, the stacking floor, and any new row are calibrated against real findings; a row that moves a large fraction of the estate is presumed

miscalibrated until argued otherwise.

6 What PAIN Relief Never Does

No entry mutates the published EPSS score or the vendor CVSS base vector; both the original and the adjusted values appear in output. No entry alters the internet-reachability determination — that is the companion exclusion-gated method’s job, and PAIN Relief merely consumes its result. No entry extends or softens a CISA KEV / BOD 26-04 remediation date; a Known Exploited Vulnerability carries its clock untouched no matter how thoroughly it is mitigated. And no entry lowers a Modified impact dimension to None; that stronger claim is a signed VEX disposition, separately governed.

7 Worked Examples

Impact — egress-deny on a Log4Shell-class finding. A CWE-917 finding scores $I = \text{High}$ on an application asset, placing it at, say, N4. The workload runs under a verified default-deny egress NetworkPolicy with no wildcard allow. Log4Shell’s remote-code-execution requires an outbound fetch of a second-stage payload; the egress policy removes the outbound leg, so the injection still fires but cannot complete. The egress-deny row lowers Modified Integrity High→Low; recomputed through the memo’s arithmetic the finding lands at N3 — a longer deadline, earned by a control the plugin read directly from the cluster. The reachability answer is unchanged: the finding is still internet-reachable; what changed is what it can accomplish.

Exploitability — blocking EDR crossing the threshold. A finding carries EPSS 0.74 (LEV, since $0.74 \geq 0.70$), is not KEV, and the floor does not apply. The workload runs blocking-mode EDR in enforcing mode, verified from its policy export, carrying $\rho = 0.85$. Then $\text{adjustedEPSS} = 0.74 \times 0.85 = 0.63 < 0.70$, and the finding moves to NLEV — the slower remediation column. Had EPSS been 0.99, the same control would leave it at 0.84, still LEV: the threshold ensures a single control relieves the marginal case, not the near-certain one. Published EPSS remains 0.74 in the record beside the adjusted value.

IN PLAIN TERMS

Same two flaws, two providers. On the hardened estate the first drops a PAIN level and the second drops onto a slower clock, purely because a scanner-adjacent tool could *see* the controls that make those flaws less dangerous here. On the unhardened estate, nothing moves. That difference is the entire point.

8 Conclusion

PAIN Relief closes the loop the PAIN memo opened. The memo re-weights a vulnerability by what is at stake on the asset; this paper re-weights it by the controls that verifiably protect the asset, using the same Environmental machinery and the same discipline of citing evidence. Impact steps down a rung when a verified control bounds the damage; exploitability crosses a threshold when verified controls make success less likely; and neither ever edits a published input, softens a KEV clock, or turns a claim into a discount. The result gives a provider what raw scanning never has: a deterministic, auditable reason that good security configuration lowers the number — and therefore a reason to build securely in the first place.