

Before the PAIN Equation: Deriving Security-Requirements Ceilings from Intended Federal Information Types

Matthew Venne
Chief Technology Officer, stackArmor

July 2026
Version 1.0

Abstract

Potential Agency Impact (PAIN) calculations require a Confidentiality, Integrity, and Availability Security Requirements vector for the affected asset. An asset archetype can provide a consistent architectural vector, but it cannot by itself answer a different question: which federal consequences are possible in a particular agency’s intended use of a Cloud Service Offering (CSO)?

This paper specifies the missing upstream derivation. It begins with the federal information types the CSO is designed and authorized to process and the information types an agency intends and is authorized to place in that offering. Confirmed overlap is mapped to the provisional objective values and special factors in NIST SP 800-60. The applicable values are aggregated independently for Confidentiality, Integrity, and Availability, preserving the full dimensional vector required by FIPS 199. That vector becomes a Security Requirements ceiling. The effective requirement for an asset is the per-objective minimum of its system-relative archetype requirement and the intended-use ceiling.

The separation matters. Archetypes remain reusable across agency deployments because they describe the role of a component inside the CSO, not every possible customer data type. The ceiling supplies the agency-specific consequence context without turning a system’s overall FIPS 199 impact label into an all-equal vector. Generic free text, file upload, or integration capability does not authorize hypothetical information types. Multi-agency assets aggregate only the definite agencies they can affect. The assessment measures adverse effect on federal operations, federal assets, and individuals; a CSP’s separate brand or corporate risk does not silently raise agency impact.

This paper ends at the effective Security Requirements vector. It does not choose a PAIN aggregation formula, word threshold, or N-rating. Those downstream policy choices belong to the PAIN method and its calibration.

Contents

1	Status, scope, and relationship to PAIN	2
2	The missing first step	3
3	Risk has a point of view	4
4	NIST SP 800-60 as provisional evidence	4
4.1	What the catalog does and does not decide	5
4.2	High-impact paths in the catalog	5

5	Preserving the FIPS 199 dimensions	6
5.1	Information-type profiles	6
5.2	System aggregation	6
5.3	Certification Class is not the vector	6
6	Intended use, not hypothetical capability	7
6.1	Two sets	7
6.2	Generic input mechanisms	7
6.3	System-generated and control information	7
7	From applied information types to a ceiling	8
7.1	One agency	8
7.2	Several agencies and affected scope	8
7.3	The vector approximation	8
8	Archetype requirements remain separate	9
8.1	Role, not software kind	9
8.2	Effective requirements	9
8.3	Why this improves reuse	10
9	Worked derivations	10
9.1	Moderate intended use on a high-role database	10
9.2	High confidentiality without H/H/H flattening	10
9.3	A shared asset across two agencies	10
9.4	A generic text field	11
9.5	Government impact and CSP reputation	11
10	Cross-CSP composition and the hyperscaler boundary	11
10.1	The knowledge problem	11
10.2	Privacy-preserving vector attestations	11
10.3	A non-normative saturation approximation	12
11	Governance and audit record	12
11.1	Minimum record	12
11.2	Change triggers	13
12	Handoff contract	13
13	Limitations and prohibited overclaims	13
14	Conclusion	14
	References	15

1 Status, scope, and relationship to PAIN

This document is informational. It does not define a FedRAMP requirement, replace an agency categorization, or make NIST SP 800-60 recommendations mandatory. It proposes a deterministic bridge from federal information categorization to the Security Requirements inputs consumed by the PAIN method.

The key words MUST, MUST NOT, SHOULD, SHOULD NOT, and MAY describe this proposed method as in RFC 2119. They do not create obligations under FedRAMP, FIPS 199, or NIST guidance.

Cloud Service Offering (CSO)

The bounded cloud service whose intended federal use is being assessed.

Information type

A category of federal information identified under NIST SP 800-60, with provisional Confidentiality, Integrity, and Availability impact recommendations.

Applied profile

The objective vector retained after an agency confirms that an information type applies and evaluates its special factors and actual use.

Archetype requirements

The uncapped CR/IR/AR vector assigned from an asset's role and consequence inside the CSO architecture.

Security-requirements ceiling

The greatest C/I/A consequence supported by the confirmed intended information use in the asset's agency scope.

Effective requirements

The per-objective minimum of the archetype requirements and the ceiling. This is the CR/IR/AR vector handed to PAIN.

SCOPE BOUNDARY

The paper stops at effective CR/IR/AR. Whether a particular effective vector can produce Minimal, Narrow, Disruptive, or Debilitating under a selected PAIN formula is a property of that downstream formula and its governed thresholds. It is not a property established by NIST SP 800-60 or FIPS 199 alone.

2 The missing first step

FedRAMP requires providers to evaluate vulnerabilities in the context of the CSO and estimate potential impact on government customers. CVSS Environmental Security Requirements supply a standard vocabulary for stating how strongly the affected environment requires confidentiality, integrity, and availability. Neither source, however, automatically provides the three values for a particular asset in a particular agency deployment.

When we originally wrote the PAIN method, we recommended asset archetypes: governed role descriptions such as system-of-record database, identity service, orchestrator, message broker, or disposable cache. We did so to replace unexplained asset-owner declarations with reusable classifications whose rationale could be reviewed and audited. In that original formulation, however, we asked the archetype's CR/IR/AR vector to carry two different kinds of information:

1. **Architectural consequence:** what role does this component perform inside the CSO, and how much can its compromise affect the offering?
2. **Federal mission consequence:** what information and government use are actually in scope for the affected agency or agencies?

That coupling was a pragmatic first step, but its limitation becomes visible when the same CSO and architectural roles are reused across agencies. Combining both questions inside every archetype makes the catalog customer-specific. A database archetype may become High for one agency, Moderate for another, and then change again when the CSO adds a customer. The resulting catalog is difficult to reuse and difficult to audit.

The refinement proposed here preserves the original governed-archetype mechanism while separating the questions. The archetype holds the first answer. A Security Requirements ceiling derived from intended federal information use holds the second. PAIN receives their per-objective intersection.

IN PLAIN TERMS

A primary database can be architecturally important even when the current federal use is only Moderate. Keep the database role stable. Derive the agency’s three-dimensional ceiling separately. The database keeps its identity, while its effective PAIN requirements reflect what is actually at stake for that agency.

3 Risk has a point of view

FIPS 199 defines impact through adverse effects on organizational operations, organizational assets, and individuals. NIST SP 800-60 applies that federal frame to information types and agency missions. FedRAMP’s PAIN vocabulary likewise asks about effects on agencies using the cloud service.

Risk is therefore not an intrinsic property of a byte sequence. The same information can carry different consequences for different stakeholders:

Stakeholder	Question	Treatment in this method
Federal agency	Would disclosure, corruption, or loss cause limited, serious, or severe or catastrophic adverse effect on the mission, federal assets, or individuals?	This is the impact measured by the ceiling and handed to PAIN.
CSP	Would the same event harm brand, market position, revenue, contractual posture, or corporate reputation?	Important enterprise risk, but not automatically federal agency impact. Include it only where it creates a documented federal consequence.
Other customer	Would the event harm a non-federal customer’s operations or obligations?	Outside federal PAIN unless that consequence also affects an agency in the assessed scope.

Table 1: Stakeholder consequences are related but not interchangeable.

A Low federal Confidentiality objective does not mean disclosure is desirable or costless. It means the agency has determined that unauthorized disclosure would have only limited adverse effect under the FIPS 199 frame. A CSP may rationally assign greater internal importance to avoiding the same disclosure. That separate judgment belongs in the CSP’s enterprise risk process rather than being smuggled into PAIN as if the government made it.

4 NIST SP 800-60 as provisional evidence

NIST SP 800-60 Volume II Revision 1 contains a catalog of management-and-support and mission-based information types. Each scored record supplies a provisional C/I/A profile, objective-level

rationale, special factors, and a recommendation. The values are starting points. The agency identifies applicable information types, evaluates actual use, applies special factors, considers aggregation and system context, and documents adjustments.

4.1 What the catalog does and does not decide

For an information type t , write its provisional profile as

$$p_t = (p_t^C, p_t^I, p_t^A), \quad p_t^O \in \{L, M, H\}.$$

After the agency evaluates actual use and special factors, the applied profile is

$$q_{a,t} = (q_{a,t}^C, q_{a,t}^I, q_{a,t}^A).$$

The method MUST preserve both values and explain every difference. A provisional High is not automatic proof that the agency’s use is High. A provisional Low does not prevent a documented special factor or aggregate from raising the applied objective. A candidate information type does not participate in the math until its applicability is confirmed.

The catalog does not cover classified information or national-security systems. Those uses remain outside this method.

4.2 High-impact paths in the catalog

A source-pinned conversion of Volume II contains 170 records: 168 scored information types and two delivery mechanisms without standalone profiles. Eleven of the 168 scored types begin with at least one provisional High objective:

ID	Information type	Provisional vector
D.2.2	Key Asset and Critical Infrastructure Protection	C:H/I:H/A:H
D.2.3	Catastrophic Defense	C:H/I:H/A:H
D.2.4	EOP Executive Functions	C:H/I:M/A:H
D.3	Intelligence Operations (Unclassified Domestic Intelligence)	C:H/I:H/A:H
D.4.1	Disaster Monitoring and Prediction	C:L/I:H/A:H
D.4.4	Emergency Response	C:L/I:H/A:H
D.5.1	Foreign Affairs	C:H/I:H/A:M
D.5.3	Global Trade	C:H/I:H/A:H
D.11.4	Space Operations	C:L/I:H/A:H
D.14.4	Health Care Delivery Services	C:L/I:H/A:L
D.17.2	Legal Defense	C:M/I:H/A:L

Table 2: The eleven information types with a provisional High objective.

Only four provisionally begin at H/H/H. Another 94 scored records contain objective-specific special-factor language that explicitly permits High under stated conditions. Sixty-three have no explicit High path in either their provisional profile or special factors.

These counts require careful interpretation. Provisional High is uncommon. Conditional High is available across a broader part of the catalog, but a conditional path does not make the information type High by itself. Common conditions concern life safety, primary mission failure, critical infrastructure, time-critical emergency or operational decisions, catastrophic market disruption, and disclosure or corruption that enables catastrophic harm. Each requires an affirmative fact about the agency’s intended use.

SCOPE BOUNDARY

The catalog supports “provisional High is rare” and “conditional High requires specific facts.” It does not support the claim that only a small fraction of information types have any conceivable High path: 105 of 168 have either a provisional or explicit conditional path. Nor does the catalog establish how many agencies actually meet those conditions.

5 Preserving the FIPS 199 dimensions

5.1 Information-type profiles

FIPS 199 expresses the security category of information as a tuple:

$$SC_t = \{(C, q_t^C), (I, q_t^I), (A, q_t^A)\}.$$

The dimensions are not interchangeable. An information type may be High for Integrity and Low for Confidentiality and Availability. That asymmetry is evidence, not noise.

5.2 System aggregation

For the set of applicable information types T in a system, FIPS 199 takes the high-water mark independently for each objective:

$$SC_{\text{system}}(C) = \max_{t \in T} q_t^C, \quad (1)$$

$$SC_{\text{system}}(I) = \max_{t \in T} q_t^I, \quad (2)$$

$$SC_{\text{system}}(A) = \max_{t \in T} q_t^A. \quad (3)$$

NIST SP 800-60 then calls for review of aggregation, critical system functionality, and other system factors, with documented adjustment where warranted.

The system’s overall impact level is the highest of the three resulting objectives:

$$\text{Impact}_{\text{system}} = \max\{SC_{\text{system}}(C), SC_{\text{system}}(I), SC_{\text{system}}(A)\}.$$

Therefore any applied High objective makes the information system High overall. A genuinely Moderate system cannot contain a High system objective. But the overall word *High* does not replace the vector. If the category is H/M/L, PAIN needs H/M/L, not H/H/H.

DERIVATION

Use the high-water mark twice, for two different outputs.

1. Take a maximum per objective to retain the system vector.
2. Take the maximum of that vector to name the overall FIPS 199 impact level.

The overall label is useful for authorization and reporting. The dimensional vector is the input needed for an objective-sensitive PAIN calculation.

5.3 Certification Class is not the vector

FedRAMP’s 2026 terminology uses Certification Classes A–D for package specifications. Current FedRAMP guidance describes those classes as levels of assurance that can support agency use cases; it does not make the Class letter a substitute for an agency’s FIPS 199 C/I/A category. In particular, the guidance states that some Class C packages may support some High security objectives.

This method therefore records Certification Class separately. Class can corroborate an intended-use assessment and selects FedRAMP processes outside this paper, but it does not flatten the ceiling. A conflict between the agency’s category, intended use, and the CSO’s authorized boundary is surfaced for review rather than resolved by silently changing the data.

6 Intended use, not hypothetical capability

6.1 Two sets

Let:

T_{CSO}

Information types the offering is designed and authorized to process, including applicable service-generated records needed to operate the offering.

T_a Information types agency a intends and is authorized to process through the offering.

The applicable set for that agency and offering is:

$$T_{a,\text{CSO}} = T_{\text{CSO}} \cap T_a. \quad (4)$$

Each entry is more than an identifier. It carries the agency-confirmed applied profile and the special-factor rationale for this use.

6.2 Generic input mechanisms

A text box can technically accept a medical narrative, criminal-justice record, tax return, trade secret, grocery list, or poem. That technical fact does not place all corresponding information types inside the intended boundary. The same is true of file uploads, object storage, email ingestion, and general APIs.

The CSO documents what it is designed and authorized to do. The agency documents what it intends and is authorized to put there. Only the confirmed overlap enters Equation (4). Prohibited information remains prohibited. Accidental or unauthorized data spillage is managed as spillage, not normalized into the ordinary categorization merely because the user interface could carry it.

This rule prevents generic product capability from becoming an automatic High designation. It also prevents the opposite error: a documented free-text, attachment, or integration workflow that is intended to carry a sensitive information type cannot be ignored just because the product has a generic label.

6.3 System-generated and control information

The intersection is not limited to agency-uploaded content. A CSO can generate audit records, identity material, security telemetry, configurations, decisions, transactions, or derived aggregates that require categorization. NIST SP 800-60 also directs reviewers to consider critical system functionality and security policy enforcement information.

The assessment therefore records:

- agency-supplied intended information;
- service-generated information;
- derived or aggregated information;
- trusted decisions and actions performed by the CSO;

- documented system-level adjustments.

These facts can alter an applied profile. They do not erase the requirement to identify the objective and rationale being changed.

7 From applied information types to a ceiling

7.1 One agency

For agency a , objective $o \in \{C, I, A\}$, and confirmed applied profile $q_{a,t}$, define:

$$E_a(o) = \max_{t \in T_{a,CSO}} q_{a,t}^o. \quad (5)$$

Apply any documented system-level adjustment after the information-type maximum. The result

$$E_a = (E_a(C), E_a(I), E_a(A))$$

is the intended-use ceiling for that agency's deployment.

If no intended information types have been confirmed, the method does not invent a low ceiling. It records an unresolved assessment and uses a conservative, explicitly provisional upper bound until the intended use is confirmed.

7.2 Several agencies and affected scope

Let A_x be the definite agencies whose information or mission can be affected through asset x . The asset-scoped ceiling is:

$$E_x(o) = \max_{a \in A_x} E_a(o). \quad (6)$$

This is deliberately asset-scoped:

- A single-tenant component receives the ceiling for its one definite agency.
- A shared control plane receives the per-objective maximum for every definite agency it can affect.
- Adding an agency to a shared component can raise, but never lower, its ceiling.
- Adding an agency elsewhere does not change a component that cannot affect that agency.

Agency identity alone does not determine an objective. An agency may operate many systems with different information types. Only its intended use of the assessed CSO participates.

7.3 The vector approximation

Some implementations already maintain only two aggregate vectors:

- an offering or System Security Objectives vector (SSO);
- an Agency Security Objectives vector (ASO).

They can approximate Equation (5) with:

$$\hat{E}_a(o) = \min(\text{SSO}(o), \text{ASO}_a(o)). \quad (7)$$

For several agencies, take the per-objective maximum of their ASO vectors before the minimum. This approximation is simple and safe when the aggregate vectors faithfully describe the same intended-use set. It can be imprecise when the CSO and agency vectors become High for unrelated information types. Two aggregate High values prove dimensional compatibility, not type-level overlap.

The preferred audit record therefore retains confirmed information-type membership. Equation (7) remains a documented fallback, not a claim that the set intersection was evaluated.

IN PLAIN TERMS

The vector shortcut asks, “How high can each side go?” The type-aware method also asks, “Are they talking about the same intended use?” Keep the shortcut where detail is unavailable, but label it honestly.

8 Archetype requirements remain separate

8.1 Role, not software kind

An archetype describes how an asset is used inside the CSO. The same software can serve different roles:

- an in-memory store used as a disposable cache;
- the same store holding durable sessions or authentication tokens;
- the same store brokering mission-driving work;
- the same engine acting as a system-of-record database.

Those roles justify different uncapped requirements even before the agency context is applied. The archetype catalog is therefore system-relative: a provider derives it from the CSO architecture, trust relationships, control functions, and consequences of component compromise.

8.2 Effective requirements

Let the archetype requirements for asset x be

$$R_x = (R_x(C), R_x(I), R_x(A)).$$

The effective requirements handed to PAIN are:

$$R_x^{\text{eff}}(o) = \min(R_x(o), E_x(o)), \quad o \in \{C, I, A\}. \quad (8)$$

The ceiling never raises an archetype. A High ceiling means only that High remains available where the asset’s architectural role supports it. Conversely, a High archetype does not manufacture a High federal consequence when the intended-use ceiling is Moderate.

Input	Example	Scope	Meaning
Archetype requirements	H/H/M	Asset in the CSO	Uncapped architectural consequence of the role
Intended-use ceiling	M/M/L	Agency or affected agency set	Maximum federal consequence supported by confirmed intended use
Effective requirements	M/M/L	Asset in that agency scope	Per-objective minimum supplied to PAIN

Table 3: The three vectors remain visible in the audit chain.

8.3 Why this improves reuse

The same CSO can retain one governed archetype catalog across several agency deployments. The customer-specific information appears in the ceiling. A role change still changes the archetype. An intended-use change still changes the ceiling. Reviewers can challenge either judgment without rebuilding both.

Cross-system trust anchors, software-delivery paths onto agency devices, and shared CSP infrastructure require particular care because their consequences may extend beyond the information stored in the assessed component. The correct response is to model their architectural role and affected scope explicitly, not to make a data-only archetype and then hide an exception.

9 Worked derivations

9.1 Moderate intended use on a high-role database

A CSO’s system-of-record database has archetype requirements H/H/M. One agency’s confirmed intended information types aggregate to M/M/L. Therefore:

$$R_{\text{db}}^{\text{eff}} = \min((H, H, M), (M, M, L)) = (M, M, L).$$

The database remains the system-of-record archetype. No customer-specific variant is created. If a later agency uses the same CSO differently, its ceiling changes without rewriting the database role.

9.2 High confidentiality without H/H/H flattening

An agency’s applied information types and special factors aggregate to H/M/L. The overall FIPS 199 impact level is High because one objective is High. The ceiling remains H/M/L.

For a secrets service with H/H/H archetype requirements:

$$R_{\text{secrets}}^{\text{eff}} = \min((H, H, H), (H, M, L)) = (H, M, L).$$

Turning the overall High label into H/H/H would discard the agency’s dimensional judgment and overstate Integrity and Availability.

9.3 A shared asset across two agencies

Agency 1 has ceiling M/M/L. Agency 2 has ceiling L/H/M. A component dedicated to Agency 1 retains M/M/L. A component dedicated to Agency 2 retains L/H/M. A shared control plane that can affect both receives:

$$E_{\text{shared}} = \max((M, M, L), (L, H, M)) = (M, H, M).$$

Multi-agency scope increases only the objectives supported by a definite affected agency. It does not automatically force H/H/H.

9.4 A generic text field

A project-management CSO provides comments and attachments. Its documented intended use covers project plans, task records, and approved supporting artifacts. The agency has prohibited health, tax, criminal-justice, and classified content.

The text field does not add every NIST information type to T_{CSO} . Confirmed project and planning types enter the ceiling. Prohibited content remains governed by data-use policy and spillage response. If the agency later authorizes a new intended record type, the type inventory and ceiling are reviewed as a change.

9.5 Government impact and CSP reputation

An information type has applied C:L because disclosure would have limited adverse effect on the agency's operations. The CSP believes any public incident involving the record would materially harm its brand. The ceiling retains C:L unless the CSP can show how that reputational event creates a greater adverse effect on the agency, federal assets, or affected individuals. The CSP can and should manage its separate corporate exposure through enterprise risk controls.

10 Cross-CSP composition and the hyperscaler boundary

10.1 The knowledge problem

An upstream CSP may support many downstream CSOs without knowing which agencies ultimately depend on each service or which information types those agencies process. Requiring a hyperscaler to collect every downstream agency customer list is unlikely to be practical and may conflict with contractual boundaries.

The missing information does not prove a low ceiling. It creates uncertainty about the dimensional aggregate.

10.2 Privacy-preserving vector attestations

A downstream CSO could share a bounded assertion without naming its agency customers or detailed information types:

- downstream CSO identifier;
- authorization or categorization boundary;
- maximum required C/I/A ceiling for the dependency;
- whether the dependency is single-customer or shared;
- evidence status, effective date, and expiry;
- attesting party and signature.

The upstream CSP can take the per-objective maximum of current assertions. This is less precise than type-level evidence but more informative than customer identities alone.

10.3 A non-normative saturation approximation

Where no downstream vector assertions exist and the upstream service supports a broad, heterogeneous population of confirmed downstream CSOs, the dimensional maximum will tend toward the upper envelope allowed by the upstream boundary. A reasonable conservative approximation is:

$$E_{\text{downstream}} \approx \begin{cases} (M, M, M), & \text{Moderate-impact upstream system,} \\ (H, H, H), & \text{High-impact upstream system.} \end{cases}$$

This is an uncertainty policy, not a NIST or FedRAMP derivation. The trigger cannot be “many” in undocumented prose. An adopter must govern the dependency-count or diversity threshold, evidence window, confidence, and review cycle. The approximation never authorizes a use outside the upstream boundary: if a downstream use actually requires High while the upstream information system is genuinely Moderate, the result is a boundary mismatch, not permission to process the use.

SCOPE BOUNDARY

This paper identifies the cross-CSP composition problem but does not standardize a saturation threshold or attestation protocol. The approximation is suitable as a disclosed conservative fallback. A complete supply-chain profile is future work.

11 Governance and audit record

A deterministic equation does not make uncertain inputs certain. A conforming assessment preserves evidence and ownership for each judgment.

11.1 Minimum record

The record SHOULD contain:

- CSO identity, boundary, designed functions, and service-generated data;
- confirmed intended and prohibited information types;
- per-agency information types, relationship (definite or target), and governing evidence;
- each provisional and applied C/I/A profile;
- special factors considered and objective-specific adjustment rationale;
- system-level aggregation or critical-function adjustments;
- the per-agency and asset-scoped ceilings;
- the overall FIPS 199 impact label, stored separately from the vector;
- archetype requirements, effective requirements, and every capped objective;
- multi-agency affected-scope basis;
- confidence, assumptions, manual-review items, effective date, and expiry;
- any vector approximation or downstream-CSP saturation approximation used.

Operator attestations remain separate from analyst or automated inferences. Confidence measures evidence quality and never lowers an objective. When two adjacent values remain credible, select the higher value provisionally and state what evidence would justify lowering it.

11.2 Change triggers

Reassess the ceiling when:

- an agency adds or removes an intended information type;
- a special factor becomes applicable or ceases to apply;
- a CSO adds a new ingestion path, generated record, trusted decision, or system function;
- tenancy or shared-control-plane scope changes;
- aggregation creates a new system-level consequence;
- the governing categorization or authorization boundary changes;
- a downstream-CSP assertion expires or the approximation threshold is crossed.

Archetype changes follow different triggers: component role, trust, privilege, dependency, or architectural consequence. Keeping the triggers separate is one of the model's principal governance benefits.

12 Handoff contract

The output of this paper is not a PAIN word or N-rating. It is an auditable, three-dimensional input:

DERIVATION

For each affected asset x :

$$R_x^{\text{eff}}(o) = \min \left(R_x(o), \max_{a \in A_x} \max_{t \in T_{\text{CSO}} \cap T_a} q_{a,t}^o \right)$$

for $o \in \{C, I, A\}$, with documented system-level adjustments applied where required.

A transport may encode the resulting vector as:

CR:M/IR:H/AR:L

or:

cr-m_ir-h_ar-l

The handoff must also identify the uncapped archetype vector, ceiling vector, ceiling source and scope, and whether any objective was capped. The downstream PAIN method then combines this effective requirements vector with vulnerability impact.

13 Limitations and prohibited overclaims

- NIST SP 800-60 profiles are provisional recommendations. They do not replace agency categorization, governing sources, or actual-use evidence.
- Classified information and national-security systems are outside the publication's scope.
- Information-type intersection does not eliminate system-level aggregation, critical-function, control-information, or extenuating-circumstance review.

- The vector approximation can conceal mismatched information types and must be labeled as an approximation.
- The hyperscaler saturation rule is conservative policy, not a measured probability or a federal standard.
- Certification Class does not replace the FIPS 199 category or the intended-use vector.
- The ceiling is a cap, not a floor. It cannot make a low-role asset High.
- This paper does not prove that any effective vector produces a particular PAIN word. Such a result depends on the downstream PAIN formula and thresholds.
- The catalog counts describe text and provisional profiles, not the prevalence of High-impact agency deployments.

14 Conclusion

The first input to an agency-specific PAIN calculation should be the federal consequence that is actually possible in the assessed use. NIST SP 800-60 supplies provisional, objective-specific evidence; FIPS 199 supplies dimensional aggregation and the overall system impact rule. Neither calls for replacing the resulting vector with one repeated High, Moderate, or Low label.

Separating the intended-use ceiling from the asset archetype makes both judgments clearer. The archetype records what the component does in the CSO. The ceiling records what the definite agency use puts at stake. Their per-objective minimum becomes the effective CR/IR/AR vector handed to PAIN. Multi-agency components aggregate only the agencies they can affect. Generic technical capability does not authorize hypothetical data. CSP corporate consequences remain important but do not silently redefine government impact.

The result is a reusable architectural classification joined to an agency-specific, source-traceable federal consequence vector. The PAIN equation can then begin with its most consequential assumption made explicit.

References

- [FIPS199] National Institute of Standards and Technology, *Standards for Security Categorization of Federal Information and Information Systems*, FIPS PUB 199, February 2004. <https://doi.org/10.6028/NIST.FIPS.199>.
- [SP800-60-1] K. Stine, R. Kissel, W. Barker, J. Fahlsing, and J. Gulick, *Guide for Mapping Types of Information and Information Systems to Security Categories, Volume I*, NIST SP 800-60 Volume I Revision 1, August 2008. <https://doi.org/10.6028/NIST.SP.800-60v1r1>.
- [SP800-60-2] K. Stine, R. Kissel, W. Barker, A. Lee, J. Fahlsing, and J. Gulick, *Guide for Mapping Types of Information and Information Systems to Security Categories, Volume II*, NIST SP 800-60 Volume II Revision 1, August 2008. <https://doi.org/10.6028/NIST.SP.800-60v2r1>.
- [CVSS31] FIRST, *Common Vulnerability Scoring System v3.1: Specification Document*, Environmental Security Requirements. <https://www.first.org/cvss/v3.1/specification-document>.
- [FEDRAMP-VER] FedRAMP, *Vulnerability Evaluation and Reporting*, Consolidated Rules for 2026, rule VER-EVA-EPA, launched June 24, 2026. <https://www.fedramp.gov/2026/reference/rev5/b/vulnerability-evaluation-and-reporting/>.
- [FEDRAMP-PAIN] FedRAMP, *Initial Outcome of RFC-0031: Updated Incident Communications Procedures*, Notice 0012, PAIN and customer-effect terminology. <https://www.fedramp.gov/notices/0012/>.
- [FEDRAMP-CLASS] FedRAMP, *Choosing a Certification Class*, Consolidated Rules for 2026. <https://www.fedramp.gov/preview/2026/providers/start/class/>.
- [PAIN] M. Venne, *A Deterministic, CVSS-Environmental Method for FedRAMP VDR/VER Vulnerability Prioritization*, stackArmor, July 2026. <https://stackarmor.github.io/rfc-fedramp-vdr/vdr-pain-cvss.pdf>.

Reproducibility record

The catalog analysis uses the source-pinned NIST SP 800-60 Volume II Revision 1 conversion distributed with `trivy-plugin-vdr-skills`. A conforming run must return:

Check	Expected result
Total catalog records	170
Scored information types	168
Non-scoring delivery mechanisms	2
Provisional profile contains at least one High	11
Provisional H/H/H	4
Additional records with explicit conditional High language	94
Unique provisional or conditional High path	105
No explicit provisional or conditional High path	63

The conditional count includes a record only when objective-specific `specialFactors` text expressly permits a High impact level. A conditional record contributes no High objective to an assessed system unless the agency confirms that the stated factor applies.